

Flow

A High Performance Spark Data Processing Library

Chris Twiner

Copyright @ 2022 - UBS AG

Table of contents

1. Flow - 0.0.1-poc6	3
1.1 Process a Graph of data processing Steps with Quality	3
1.2 Flow and FlowT?	3
1.3 Serialisation	4
1.4 Timeouts and Error handling	4
2. About	6
2.1 Why Flow?	6
3. Scala Docs	7

1. Flow - 0.0.1-poc6

Average			
Statement	96.48 %	Branch	90.74 %

1.1 Process a Graph of data processing Steps with Quality

Write auditable DQ and transformation rules using simple SQL or create re-usable functions via SQL Lambdas and manage them through a graph of Steps.

Your Flows can be provided DataFrames from either a Quality DataFrameLoader, catalog or even custom loading logic via extension of FlowT.

-  Initial Release

Entire Flows are lazily evaluated but support caching of a Steps results for interim performance benefits.

Each Step in a Flow has an associated Runner and, unless configured not to, keeps an audit of all the results of the previous Steps in the graph, either Quality:

1. engine,
2. collector,
3. folder,
4. DQ or
5. a custom implementation

The resulting Column then has a ResultApproach applied to it, either:

- AsIs - the audit and result column is added to the DataFrame,
- MergeFields - any nested result fields are merged with the DataFrame,
- OutputFieldsOnly - only the audit column and nested result columns are output,
- StarOnly - the audit column is dropped and only nested result columns are output,
- OutFieldOnly - only the result column remains or
- a custom implementation

1.2 Flow and FlowT?

The default Flow is a type alias for:

```
type Flow = FlowT[CombinedRuleSuiteRows, RuleSuite]
```

The first parameter is for the type of Flow / "group level" rules, these may be accessed as nested runners within transformations, the second parameter is the type of RuleSuiteTypeParam parameter to use, which are passed to the Runner functions.

Flow supports both GroupRuleId and RuleSuite RuleSuiteTypeParam in Steps *and* also Id's at group level, with the exact types used following through to serialisation.

1.2.1 What is returned?

Flow's Steps are processed as graphs of Futures, each effectively returning a DataFrame, which can be optionally cached, forming the inputs to the next dependent Step. These Futures are then sequenced to return a FlowResult:

```
case class FlowResult[T: RuleSuiteTypeParam](stepResults: Map[String, StepResultType[T]],
                                             duration: Duration) extends Serializable
```

wherein each `StepResultType` is either a `StepResult` or a `StepException`, by default any return is only `StepResults`, use the `tolerant` parameter to decide how to handle failure.

```
case class StepResult[RP: RuleSuiteTypeParam](step: Step[RP], output: DataFrame, timings: StepTimings)
```

`StepResult` itself provides the resulting `DataFrame`, as the `Flow` can have multiple final leaf nodes the caller must know what to do with them, for example custom `ResultApproaches` may save the `DataFrame` to disk, or the `Flow.run` caller has custom knowledge of the flows.

1.3 Serialisation

`Flow` supports both individual datasets for a more normalised flow management approach and a `FullFlow` approach allowing individual flows to be serialised via json, for example. Although the default behaviour of both `Flow`'s and serialisation is group level `CombinedRuleSuites` and `RuleSuites`, the types can be converted to group level rules and `GroupRuleId`'s via the `convertToIds` function.

As the `Id`'s and `Rules` can be sourced from different locations in a `Flow` there are a number of strategies that can be provided to serialisation functions:

- `RuleSuiteFromDataset`
- `RuleSuiteFromFlows`
- `IdFromFlows`
- `IdFromDatasets`

these customise the processing of configuration data.

Serialising the default `Flow` to a `FullFlow` will yield a `RuleSuite` at each `Step`, doing so with `GroupId`'s as the `Step` type just has the `ruleSuiteId` and `ruleSuiteVersion` saved.

`Flow` makes use of the `sparkutils` fork of `Frameless` to use `AgnosticEncoding` for Spark 4+ for correct type handling, this also allows deciding on which type of `FlowT` should be used to be abstracted and combined by `Frameless`' encoder derivation approach.

1.4 Timeouts and Error handling

1.4.1 Timeouts

By default, a `Flow` will run until normal completion, but in many cases `Flows` will be part of business processes with tight SLAs where detecting errand processes early is important. As such there are two layers of timeouts..

1. Flow level - configured via `FlowT` parameters and `FlowRow` "duration" fields.
2. Step level - configured via a `Step`'s `stepTimeout` options

Step level timeouts can operate independently, allowing a given problematic step to be configured with a timeout that then applies to the whole `Step` sub-graph.

When a timeout is reached a `StepException` is returned for a `Step` instead of a `StepResult`.

1.4.2 Exception / Error handling

When calling `Flow.run` the default value of the "tolerant" parameter is 'false' e.g.:

```
val ires = flow.run(sparkSession)
```

In this default behaviour **any** exception, whether in Spark or via `stepTimeout` options, will trigger the run call to throw.

In contrast:

```
val ires = flow.run(sparkSession, tolerant = true)
```

will return a mix of responses for each Step. You can use either pattern matching or the `StepResultType.fold` function to process failures and successes.

Last update: May 3, 2026 16:15:28

Created: May 3, 2026 16:15:28

2. About

2.1 Why Flow?

In short, the need for a lightweight configurable and auditable transformation approach integrated within a Spark application wasn't being met by other identified solutions. Many full-fledged products required additional infra, or required using Python instead of configuration through data suitable for audit and end-user maintenance.

Flow aims to fill the 'in process' data processing library gap configurable via simple sql snippets and some json.

Last update: May 3, 2026 16:15:28

Created: May 3, 2026 16:15:28

3. Scala Docs

[flow_api](#) docs.

The server side interface can be found [here](#).

Last update: May 3, 2026 16:15:28

Created: May 3, 2026 16:15:28